



Sociotechnical AI Governance: Challenges and Opportunities for HCI

K. J. Kevin Feng*
Human Centered Design &
Engineering
University of Washington
Seattle, Washington, USA
kjfeng@uw.edu

Rock Yuren Pang*
Paul G. Allen School of Computer
Science & Engineering
University of Washington
Seattle, Washington, USA
ypang2@cs.washington.edu

Tzu-Sheng Kuo
Human-Computer Interaction
Institute
Carnegie Mellon University
Pittsburgh, Pennsylvania, USA
tzushenk@cs.cmu.edu

Amy Winecoff
Center for Democracy & Technology
Washington, District of Columbia
USA
awinecoff@cdt.org

Emily Tseng
Microsoft Research
New York, New York, USA
emilytseng@microsoft.com

David Gray Widder
Cornell Tech
New York, New York, USA
david.g.widder@gmail.com

Harini Suresh
Brown University
Providence, Rhode Island, USA
harinisuresh824@gmail.com

Katharina Reinecke
Paul G. Allen School of Computer
Science & Engineering
University of Washington
Seattle, Washington, USA
reinecke@cs.washington.edu

Amy X. Zhang
Paul G. Allen School of Computer
Science & Engineering
University of Washington
Seattle, Washington, USA
axz@cs.uw.edu

Abstract

Rapid advancements in and adoption of frontier AI systems have amplified the need for AI governance measures across the public sector, academia, and industry. Prior work in technical AI governance has proposed agendas for governing technical components in AI development, such as data, models, and compute. However, recent calls for more *sociotechnical* approaches recognize the critical role of social infrastructures surrounding technical ones in shaping governance decisions and efforts. While scholars and practitioners have advocated for sociotechnical AI governance, concrete research directions in this area are only beginning to emerge. This workshop aims to gather the expertise of researchers in HCI and adjacent disciplines to chart promising paths forward for sociotechnical AI governance. To make problems in this area more tangible, we outline four core governance challenges for contributions: **anticipating** high-priority risks to address with governance, **identifying** where to focus governance efforts and who should lead those efforts, **designing** appropriate interventions and tools to implement governance actions in practice, and **evaluating** the effectiveness of these interventions and tools in context. Through papers, panel discussions, keynotes, and collaborative drafting of a research agenda, this

workshop will build community and empower actionable efforts to tackle AI governance through a sociotechnical lens.

CCS Concepts

• **Computing methodologies** → **Artificial intelligence**; • **Human-centered computing** → **Collaborative and social computing theory, concepts and paradigms**; **Human computer interaction (HCI)**; • **Social and professional topics** → **Computing / technology policy**.

Keywords

AI governance, sociotechnical methods, responsible AI, human-AI interaction

ACM Reference Format:

K. J. Kevin Feng, Rock Yuren Pang, Tzu-Sheng Kuo, Amy Winecoff, Emily Tseng, David Gray Widder, Harini Suresh, Katharina Reinecke, and Amy X. Zhang. 2025. Sociotechnical AI Governance: Challenges and Opportunities for HCI. In *Extended Abstracts of the CHI Conference on Human Factors in Computing Systems (CHI EA '25)*, April 26–May 01, 2025, Yokohama, Japan. ACM, New York, NY, USA, 6 pages. <https://doi.org/10.1145/3706599.3706747>

1 Motivation

AI governance is broadly defined as the “processes and structures through which decisions related to AI are made, implemented, and enforced” [17]. As the development and adoption of frontier AI systems (e.g., OpenAI’s o1, Anthropic’s Claude, Google’s Gemini) accelerate, researchers and practitioners across academia, the public sector, and industry are increasingly prioritizing AI governance initiatives to ensure AI can be safely and responsibly integrated into society [2–4, 9, 11, 13, 15]. In the United States, for example, major AI labs have voluntarily committed to the White House’s

*Both authors contributed equally.

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for third-party components of this work must be honored. For all other uses, contact the owner/author(s).

CHI EA '25, Yokohama, Japan

© 2025 Copyright held by the owner/author(s).

ACM ISBN 979-8-4007-1395-8/25/04

<https://doi.org/10.1145/3706599.3706747>

Executive Order for “Safe, Secure, and Trustworthy Development and Use of Artificial Intelligence” [13] through measures such as early warning evaluations [11]. Both the European Union (via the AI Act) and South Korea (via the AI Basic Act) have introduced comprehensive regulatory frameworks to address increasing safety, ethical, and societal concerns of AI while fostering innovation [15, 20]. Anthropic also introduced its Responsible Scaling Policy [3] to help manage safety, security, and operational standards appropriate to model capability. As model capabilities inevitably improve, the rigor and scope of governance efforts from all stakeholders will need to keep pace.

Prior work in *technical* AI governance has highlighted the need for targeted technical research and expertise in governance [6, 18], as well as identified open research questions for governing technical components of AI, including data, algorithms, compute, and deployment infrastructure [17]. While this technical approach is valuable, it is by no means sufficient on its own—scholars in HCI and beyond have long advocated for *sociotechnical* approaches to technology development. That is, we must understand and account for social factors (e.g., relevant stakeholders, workflows, collaboration dynamics, cultures) and alongside technical components, such that two are viewed as one coherent system [1, 8, 12, 14, 19, 21]. However, despite increasing calls for applying sociotechnical methods to AI governance [5, 7, 10, 16], concrete research agendas have yet to emerge. With its long history of sociotechnical work and a wide breadth of interdisciplinary expertise, we believe the CHI community is uniquely positioned to pioneer impactful research and policy contributions in this area.

We thus propose a one-day workshop at CHI 2025 to bring together researchers in HCI, human-AI interaction, and adjacent areas to build community and collaboratively draft a research agenda for sociotechnical AI governance. Importantly, We encourage stakeholders from international backgrounds to discuss AI regulation across diverse cultural and regulatory contexts. To better focus efforts, we propose four core governance challenges that will guide workshop submissions, discussions, and outcomes. These challenges are drawn from prior work from Dean et al. identifying axes of sociotechnical inquiry in AI research [8], and adapted to more directly tackle governance. The four challenges are: **anticipating** high-priority risks to address with governance, **identifying** where to focus governance efforts and who should participate in and lead those efforts, **designing** appropriate interventions and tools to implement governance actions in practice, and **evaluating** the effectiveness of these interventions and tools in context. Specific questions that attendees may tackle in our workshop include—but are certainly not limited to—the following:

- *Anticipating*: How can we identify critical risks associated with AI systems with sociotechnical methods? What are practical guidelines to anticipate and prioritize the AI risks for governance?
- *Identifying*: Which stakeholders should lead AI governance efforts under what contexts, and how can we ensure diverse representation? What role should HCI researchers play in shaping governance efforts?
- *Designing*: What processes and interactive tools can be designed to support sociotechnical AI governance? How can

we design them to be feasibly implemented and used in practice?

- *Evaluating*: How do we evaluate the success of governance efforts, in both the short and long term? How can we situate evaluation in the appropriate sociotechnical context(s)?

Our workshop aims to satisfy three objectives. First, we wish to *connect researchers across HCI, AI, the social sciences, law, and policy experts*. Sociotechnical work relies on interdisciplinary concepts and methods, and we believe deeper conversations between the HCI community and governance experts less familiar with HCI can greatly enrich efforts in this space. Second, we wish to *develop a repository of goals, resources, and research questions to share publicly*. The outcomes from the workshop will be documented and shared with all attendees, and attendees will also be invited to join a collaborative effort to assemble the repository. Finally, we wish to *build a vibrant community for sociotechnical AI governance research and practice*. We envision this workshop to seed knowledge sharing and new collaborations that drive forward impactful research and practical solutions for AI governance’s evolving challenges.

2 Organizers

K. J. Kevin Feng is a Ph.D. student in the department of Human Centered Design & Engineering at the University of Washington. His research contributes novel processes and interactive systems that enable collective input for governing AI behavior, often drawing inspiration from methods in human-centered design. He specifically focuses on engaging key stakeholders who may typically be excluded from AI development, including user experience designers, attorneys, and mental health clinicians. His work has appeared in venues including CHI, CSCW, FAccT, and DIS, and has been recognized with an OpenAI Democratic Inputs to AI grant and a UW Herbold Fellowship.

Rock Yuren Pang is a Ph.D. student at the Paul G. Allen School of Computer Science & Engineering at the University of Washington. His research aligns the goals of responsible computing with the realities of computing research practices, especially in the early stages of designing technologies. He designs systems and sociotechnical approaches to understand, incentivize, and facilitate the anticipation of consequences in computing practice. His research explores AI governance approaches that benefit developers and researchers in complex organizational structures. His work has been supported by the IBM Ph.D. Fellowship and covered by the Montreal AI Ethics Institute, UW CSE, and Cornell News.

Tzu-Sheng Kuo is a PhD student in the Human-Computer Interaction Institute at Carnegie Mellon University. His research focuses on transforming the current top-down and centralized approach to AI development, towards more community-driven, deliberative, and democratic processes. To achieve this vision, he develops systems and processes that empower individuals and communities impacted by AI to actively shape its design, evaluation, and governance. His research has received Best Paper and Honorable Mention Awards at top Human-Computer Interaction conferences, including CHI and UIST.

Amy A. Winecoff is a Senior Technologist in the AI Governance Lab at the Center for Democracy & Technology (CDT), where

she focuses on developing technically-informed solutions for effective AI regulation and governance, with the goal of advancing the interests of those impacted by AI systems. Her research emphasizes building the foundations for robust governance, particularly in the areas of AI documentation and measurement. Her work has been featured in academic venues like RecSys, CHI, and First Monday, as well as in policy-focused publications through CDT and Tech Policy Press. She has also served as a responsible technology advisor for startup accelerators, and in her prior roles as a data scientist, she developed and deployed production recommendation systems for e-commerce companies.

Emily Tseng is a postdoctoral scholar at Microsoft Research with the Social Media Collective. Her research explores how digital technologies mediate harm, how to intervene, and what it means to do so. Emily's work has featured at top-tier venues in HCI (CHI, CSCW), responsible computing (FAccT), and computer security and privacy (USENIX, Oakland), including several Best Paper distinctions. She will next join the University of Washington as an Assistant Professor in Human-Centered Design and Engineering.

David Gray Widder studies how people creating "Artificial Intelligence" systems think about the downstream harms their systems make possible, and the wider cultural, political, and economic logics which shape these thoughts. He is a Postdoctoral Fellow at the Digital Life Initiative at Cornell Tech, and earned his PhD from the School of Computer Science at Carnegie Mellon University. He has previously conducted research at Intel Labs, Microsoft Research, and NASA's Jet Propulsion Laboratory. His recent research has been accepted to FAccT, Nature, CSCW, and Big Data and Society; and covered in media outlets including Wired, the Associated Press, and the New York Times.

Harini Suresh is an Assistant Professor of Computer Science at Brown University, where she is also affiliated with the Department of Science, Technology & Society (STS) and the Center for Technological Responsibility, Reimagination, and Redesign (CNTR) at the Data Science Institute. She runs the Data in Society Collective (DISCO Lab), which strives to support more grassroots and participatory approaches to AI development. As part of the cross-institutional Counterdata Network, the lab contributes to understanding the data practices of 40+ civil society activist groups who collect missing data on human rights abuses, and to the co-design of AI-based tools to support their work. Her work has contributed to research in HCI, AI & ML, and ethical technology, appearing in venues including FAccT, CHI, IUI, and EAAMO.

Katharina Reinecke is a Professor in the Paul G. Allen School of Computer Science & Engineering, University of Washington. She researches how people's interaction with digital technology varies depending on cultural, geographic, or demographic background, and how technology can be biased against people who are unlike the small groups of people who created it. Her lab has developed systems that make technology more suitable for diverse user groups and that can help designers and developers anticipate unintended consequences of technology.

Amy X. Zhang is an Assistant Professor in the Allen School of Computer Science & Engineering at University of Washington, where she leads the Social Futures Lab, dedicated to reimagining social and collaborative systems to empower people and improve society. Her work in governance has focused on the development

of novel software tools and workflows to support broader participation, improved consistency, and greater legitimacy, drawing from theories of offline governance. She has published in venues such as CHI, CSCW, UIST, and FAccT.

3 Plans to Publish Workshop Proceedings

The workshop information will be shared via a public website, including the call for participation, the program, the schedule, the organizers, the program committee, and the accepted papers. The URL of our website is expected to be chi-staig.github.io. We will invite workshop candidates to submit position papers and/or short research papers of up to four pages (excluding references). We encourage submissions from diverse sectors—such as the public sector, academia, and industry—to expose attendees to a cross-disciplinary view of AI governance. Based on these submissions, we will select workshop attendees with the aim of covering diverse domains of expertise, perspectives, and technology backgrounds. We will share the selected workshop papers among all attendees before the week of the workshop.

4 Hybrid Format, Asynchronous Materials and Accessibility

4.1 Hybrid Format

We are planning to organize a hybrid workshop, inclusive to in-person and remote attendees. Our expected number of in-person and virtual attendees is around 30 and 15, respectively, and we specifically encourage researchers from a diverse range of disciplines to participate, including governance researchers who may not be regular attendees of CHI. Additionally, some may hesitate to attend in-person due to personal, health, or environmental concerns. We have decided to plan for a hybrid workshop for these reasons.

We will use Zoom to broadcast the opening and closing remarks, keynote, and panel in the workshop to those attending virtually. For the poster and ideation sessions, we will invite those attending online to a virtual canvas (e.g., Figma, Miro) with multiple "rooms." This canvas will 1) facilitate interaction between online and offline participants, and 2) allow for easy preservation of questions, ideas, and feedback throughout the workshop. A main page will contain all posters, which organizers will upload before the start of the workshop.

Each attendee (or group of attendees if they are co-authors) will have a separate page that contains their poster. This page serves as their personal virtual poster stand. Virtual attendees can visit others' pages and add comments, and can also do the same for copies of in-person posters on the main page. Simultaneously, virtual attendees can also present and discuss posters over Zoom. We will set up computers around the room during poster sessions to encourage in-person attendees to engage with the virtual poster stands and leave feedback. During the ideation sessions, virtual attendees will be divided into breakout rooms and continue using comments and annotations on the virtual canvas to supplement synchronous virtual discussions. For in-person roundtable discussions, attendees will also write their discussion outcomes in the shared virtual canvas, which allows virtual and in-person participants to easily view and engage with each others' work. During the

end-of-day cross-cutting discussion, in-person and online groups will have a chance to share their outcomes.

4.2 Asynchronous Materials

We will create a public workshop website for distributing information and marketing, including the workshop description, a call for participation, the schedule, and information about keynote speaker and panelists. We will also circulate accepted papers before the workshop to help participants become familiar with the paper topics. In addition, we will communicate updates via emails and a group Slack/Discord channel. The channel can also be the platform for participants to network and share similar research interests before, during, and after the workshop. The organizers will maintain the channel after the workshop and propose follow-up events if there is sufficient interest. We will asynchronously share papers, posters, out discussion outcomes from the workshop via our website for public access. Recordings of the keynote and panel discussions will also be shared on our website, as well as in the attendees' Slack/Discord channel for convenient access after the workshop.

4.3 Accessibility

We plan to send a pre-workshop survey to attendees to anticipate accessibility needs. Based on the responses, we will work with conference organizers to secure the necessary video, audio, and mobility devices to meet those needs, such as listening devices. We will instruct and remind authors to comply with SIGCHI's guide to accessible submissions¹, as well as ensure that content on our website complies with modern web accessibility standards². We will enable real-time captions for all activities involving Zoom. For in-person attendees, we will provide accessible seating.

5 Workshop Activities

Our proposed agenda for the workshop can be found in Table 1.

6 Post-Workshop Plans

Immediately following the workshop, we plan to continue conversations with a post-workshop dinner. Upon collecting attendees' consent, we will gather and distribute emails of attendees so they can stay in touch if desired. We will also invite all attendees to collaborate on a public repository of goals, research questions, and resources for sociotechnical AI governance, which we will release on our website as well as potentially a report that can be posted to a preprint server such as arXiv. We will also plan for a virtual check-in after 3 months for attendees and organizers to finalize the repository, share new updates on their governance-related projects, and discuss any new developments in the governance space.

7 Call for Participation

As AI systems become increasingly powerful and pervasive, there is an urgent need to design effective AI governance measures around both technical and social factors. Sociotechnical AI governance recognizes that AI's real-world impacts are always a product of both technical capabilities and broader social factors including stakeholders, organization structures, power dynamics, and cultural norms.

To explore this important emerging topic, we invite authors from across academia, industry, the legal domain, and the public sector to submit papers to the  **STAIG@CHI'25: Sociotechnical AI Governance**.³

This workshop aims to build community and collaboratively draft a research agenda for sociotechnical AI governance. In particular, we outline four governance challenges for authors to consider: **anticipating** high-priority risks to address with governance, **identifying** where to focus governance efforts and who should participate in and lead those efforts, **designing** appropriate interventions and tools to implement governance actions in practice, and **evaluating** the effectiveness of these interventions and tools in context.

Topics authors may choose to tackle include, but not limited to:

- Theoretical and empirical understanding of stakeholders' needs and goals in AI governance.
- Novel interactive tools and interventions for collaborative governance.
- Case studies of governance in various sociotechnical scenarios.
- Evaluation methods for governance measures in practice.

Submission requirements: Submitted papers should be up to four (4) pages in the ACM single-column format, excluding references. Submissions will go through a peer review review process based on quality. Each paper will receive two high-quality reviews from experts. We will advertise the accepted papers on our workshop website. At least one author must attend the workshop to present the paper at the workshop.

References

- [1] Mark S Ackerman. 2000. The intellectual challenge of CSCW: the gap between social requirements and technical feasibility. *Human-Computer Interaction* 15, 2-3 (2000), 179–203.
- [2] Markus Anderljung, Joslyn Barnhart, Anton Korinek, Jade Leung, Cullen O'Keefe, Jess Whittlestone, Shahar Avin, Miles Brundage, Justin Bullock, Duncan Cass-Beggs, et al. 2023. Frontier AI regulation: Managing emerging risks to public safety. *arXiv preprint arXiv:2307.03718* (2023).
- [3] Anthropic. 2023. Anthropic's Responsible Scaling Policy. <https://www.anthropic.com/news/anthropics-responsible-scaling-policy>.
- [4] ANTHONY M Barrett, Jessica Newman, Brandie Nonnecke, Dan Hendrycks, EVAN R Murphy, and Krystal Jackson. 2023. AI risk-management standards profile for general-purpose AI systems (GPAIS) and foundation models. *Center for Long-Term Cybersecurity, UC Berkeley*. <https://perma.cc/8W6P-2UUK> (2023).
- [5] Miranda Bogen and Amy Winecoff. 2024. Applying a sociotechnical approaches to AI governance in practice. *Center for Democracy and Technology* (2024).
- [6] Samuel R Bowman, Jeeyoon Hyun, Ethan Perez, Edwin Chen, Craig Pettit, Scott Heiner, Kamile Lukošūtė, Amanda Askell, Andy Jones, Anna Chen, et al. 2022. Measuring progress on scalable oversight for large language models. *arXiv preprint arXiv:2211.03540* (2022).
- [7] Brian Chen and Jacob Metcalf. 2024. Explainer: A Sociotechnical Approach to AI Policy. *Data & Society* (2024).
- [8] Sarah Dean, Thomas Krendl Gilbert, Nathan Lambert, and Tom Zick. 2021. Axes for sociotechnical inquiry in AI research. *IEEE Transactions on Technology and Society* 2, 2 (2021), 62–70.
- [9] Kimberly Do, Rock Yuren Pang, Jiachen Jiang, and Katharina Reinecke. 2023. "That's important, but...": How Computer Science Researchers Anticipate Unintended Consequences of Their Research Innovations (*CHI '23*). Association for Computing Machinery, New York, NY, USA, Article 602, 16 pages. <https://doi.org/10.1145/3544548.3581347>
- [10] Roel Dobbe and Anouk Wolters. 2024. Toward Sociotechnical AI: Mapping Vulnerabilities for Machine Learning in Context. *Minds and Machines* 34, 2 (2024), 1–51.

¹<https://sigchi.org/resources/guides-for-authors/accessibility/>

²<https://www.w3.org/WAI/standards-guidelines/>

³chi-staig.github.io

- [11] Anca Dragan, Helen King, and Allan Dafoe. 2024. Introducing the Frontier Safety Framework. <https://deepmind.google/discover/blog/introducing-the-frontier-safety-framework/>.
- [12] Upol Ehsan, Koustuv Saha, Munmun De Choudhury, and Mark O Riedl. 2023. Charting the sociotechnical gap in explainable ai: A framework to address the gap in xai. *Proceedings of the ACM on human-computer interaction* 7, CSCW1 (2023), 1–32.
- [13] The White House. 2023. Executive Order on the Safe, Secure, and Trustworthy Development and Use of Artificial Intelligence. <https://www.whitehouse.gov/briefing-room/presidential-actions/2023/10/30/executive-order-on-the-safe-secure-and-trustworthy-development-and-use-of-artificial-intelligence/>.
- [14] Carsten S Østerlund, Pernille Bjørn, Paul Dourish, Richard Harper, and Daniela K Rosner. 2015. Sociomateriality and design. In *Proceedings of the 18th ACM Conference Companion on Computer Supported Cooperative Work & Social Computing*. 126–130.
- [15] European Parliament. 2023. EU AI Act: first regulation on artificial intelligence. <https://www.europarl.europa.eu/topics/en/article/20230601STO93804/eu-ai-act-first-regulation-on-artificial-intelligence>.
- [16] Inioluwa Deborah Raji and Roel Dobbe. 2023. Concrete problems in AI safety, revisited. *arXiv preprint arXiv:2401.10899* (2023).
- [17] Anka Reuel, Ben Bucknall, Stephen Casper, Tim Fist, Lisa Soder, Onni Aarne, Lewis Hammond, Lujain Ibrahim, Alan Chan, Peter Wills, et al. 2024. Open problems in technical ai governance. *arXiv preprint arXiv:2407.14981* (2024).
- [18] Anka Reuel, Lisa Soder, Ben Bucknall, and Trond Arne Undheim. 2024. Position Paper: Technical Research and Talent is Needed for Effective AI Governance. *arXiv preprint arXiv:2406.06987* (2024).
- [19] Andrew D Selbst, Danah Boyd, Sorelle A Friedler, Suresh Venkatasubramanian, and Janet Vertesi. 2019. Fairness and abstraction in sociotechnical systems. In *Proceedings of the conference on fairness, accountability, and transparency*. 59–68.
- [20] Margareth Theresia. 2024. Newly enacted law sets basis for nat'l development of AI. <https://www.korea.net/NewsFocus/policies/view?articleId=264071>.
- [21] Hubert D Zajac, Dana Li, Xiang Dai, Jonathan F Carlsen, Finn Kensing, and Tariq O Andersen. 2023. Clinician-facing AI in the Wild: Taking Stock of the Sociotechnical Challenges and Opportunities for HCI. *ACM Transactions on Computer-Human Interaction* 30, 2 (2023), 1–39.

Time	Activity
9:00–9:15	Opening remarks. The organizers will introduce themselves and the workshop goals, and review a plan of the day.
9:15–10:00	Opening panel. This panel discussion will attempt to define “sociotechnical AI governance.” Panel members will be from diverse sectors, within and outside of the CHI community. A secondary aim of the panel is to pinpoint fertile areas of overlap between broader governance efforts and the theories, methods, and techniques well-known to HCI. There will be 30 minutes of organizer-moderated discussion, followed by 15 minutes of discussion from audience questions.
10:00–10:15	Coffee break. During this time, half of the attendees with accepted papers (assigned by organizers before the workshop) will put up their posters in preparation for the poster session.
10:15–11:15	Morning poster session. Attendees assigned to the morning poster session will present their posters in a casual and conversational setting. Presenters will be encouraged to prepare a question for poster viewers, which will be further explored in subsequent ideation sessions.
11:15–12:00	Morning ideation session. Those who just presented posters will co-lead a roundtable discussion at a table that corresponds to the core AI governance challenge of their choice (typically, the challenge their work most directly addresses). The goal of the discussion is to identify concrete open research questions within a particular governance challenge and brainstorm approaches to tackle them. Each roundtable will write their discussion outcomes on a shared virtual canvas.
12:00–13:30	Lunch break.
13:30–14:15	Keynote speaker. A seasoned researcher in the field will give a 30-minute talk on a topic of choice related to sociotechnical AI governance. This talk is followed by a 15-minute audience Q&A.
14:15–14:30	Coffee break. Attendees assigned to the afternoon poster session will prepare their posters.
14:30–15:30	Afternoon poster session. Attendees assigned to the afternoon poster session will present their posters in a casual and conversational setting.
15:30–16:15	Afternoon ideation session. Those who just presented posters will co-lead a roundtable discussion at a table that corresponds to a AI core governance challenge of their choice. Each roundtable will write their discussion outcomes on a shared virtual canvas, adding onto the content from the morning ideation sessions.
16:15–16:45	Cross-cutting discussion. We will provide a chance for everyone to view all shared canvases produced. Attendees will read each canvas and post comments/feedback for each in the form of virtual sticky notes. Virtual and in-person attendee groups will then each share their discussion outcomes.
16:45–17:00	Closing remarks. Organizers will deliver concluding comments and invite all attendees to participate in the assembling of discussion outcomes for all 4 core AI governance challenges into a publicly shareable resource. We will also take group photos and encourage attendees to exchange contact information.
18:00–20:00	Offsite dinner and networking.

Table 1: Proposed workshop agenda.